

# NLP TP3 REPORT

Model Improvement

SEHILI CHAIMA  
GROUP 01

2024

The previous models that achieved approximately **66%** accuracy on the test set employed **TF-IDF** and **one-hot encoding** vectorization techniques, so I chose to use the TF-IDF vectorization technique and do the study on multiple machine learning algorithms and deep neural network architectures.

## PREVIOUSLY USED ARCHITECTURE:

A simple **MLP** consists of a **single hidden layer** with **50 neurons**. It utilizes the **ReLU** activation function, employs the **Adam** solver for optimization,

```
tfidf_classifier = MLPClassifier(hidden_layer_sizes=(50),  
                                max_iter=300, activation = 'relu',  
                                solver='adam',  
                                random_state=0)
```

```
MLPClassifier  
MLPClassifier(hidden_layer_sizes=50, max_iter=300, random_state=0)
```

## PREVIOUS RESULTS :

Classification Report for TF-IDF Representation:

	precision	recall	f1-score	support
0	0.73	0.71	0.72	2370
1	0.73	0.70	0.71	1690
2	0.69	0.70	0.69	1814
micro avg	0.71	0.70	0.71	5874
macro avg	0.71	0.70	0.71	5874
weighted avg	0.71	0.70	0.71	5874
samples avg	0.74	0.70	0.69	5874

Accuracy for TF-IDF Representation: 0.665645216207014

# TOKENIZATION TECHNIQUE UPDATE: RULE BASED TOKENIZATION ( TREEBANK TOKENIZER )

[never, occur, fumbly, might, mere, mis]

After the change of the tokenization method from subword tokenization using sentence piece to rule based tokenization using treebank and holding the same MLP architecture, i noticed an improvement in the model's accuracy to 75%.

## Results after changing the tokenization technique and keeping the same MLP architecture :

Classification Report for TF-IDF Representation:				
	precision	recall	f1-score	support
0	0.76	0.78	0.77	1580
1	0.78	0.72	0.75	1127
2	0.72	0.75	0.73	1209
accuracy			0.75	3916
macro avg	0.75	0.75	0.75	3916
weighted avg	0.75	0.75	0.75	3916
Accuracy for TF-IDF Representation: 0.7502553626149132				

Keeping the rule based tokenization technique and move forward to building our models ...

Splitting the data to x\_train, x\_test, y\_train, y\_test and showing the corresponding shape of each:

```
X shape: (19579, 1)
X train shape: (15663, 1)
X test shape: (3916, 1)
Y shape: (19579,)
Y train shape: (15663,)
Y test shape: (3916,)
```

# MACHINE LEARNING ALGORITHMS:

## 1- LOGISTIC REGRESSION:

```
from sklearn import linear_model
classes=np.unique(y_train)
lr = linear_model.LogisticRegression(multi_class='multinomial',
                                     solver='sag',
                                     class_weight={c: 1 for c in classes})
lr.fit(x_train_tfidf, y_train)
```

✓ 4m 3.5s

```
LogisticRegression
LogisticRegression(class_weight={0: 1, 1: 1, 2: 1}, multi_class='multinomial',
                    solver='sag')
```

Classification report for N-1st representation:

	precision	recall	f1-score	support
0	0.78	0.87	0.82	1580
1	0.83	0.78	0.80	1127
2	0.84	0.77	0.80	1209
accuracy			0.81	3916
macro avg	0.82	0.80	0.81	3916
weighted avg	0.81	0.81	0.81	3916

The classification report,  
accuracy and log loss:

ACCURACY:  
0.8097548518896833

LOG LOSS LR:  
0.5421346903824495

## 2- SIMPLE NAIVE BAYES:

```
from sklearn.naive_bayes import MultinomialNB
nb = MultinomialNB()
nb.fit(x_train_tfidf, y_train)
```

✓ 0.8s

```
▼ MultinomialNB
MultinomialNB()
```

	precision	recall	f1-score	support
0	0.76	0.88	0.82	1580
1	0.90	0.73	0.81	1127
2	0.83	0.81	0.82	1209
accuracy			0.82	3916
macro avg	0.83	0.81	0.81	3916
weighted avg	0.82	0.82	0.82	3916

The classification report,  
accuracy and log loss:

ACCURACY:  
0.8153728294177732

LOG LOSS NB:  
0.5849931054230106

- VERY FAST TRAINING  
AND PREDICTION  
TIME!

# 3- SVM (SUPPORT VECTOR MACHINES):

```
from sklearn.svm import SVC
svm= SVC(kernel='linear',probability=True)
svm.fit(x_train_tfidf,y_train)
```

```
SVC
SVC(kernel='linear', probability=True)
```

ACCURACY:  
0.8010725229826353

LOG LOSS SVM:  
0.47319472036205884

- VERY LONG TRAINING AND PREDICTION TIME.

# 4- XGBOOST:

```
import xgboost as XGB
params = {
    'objective':'binary:logistic',
    'max_depth': 4,
    'alpha': 10,
    'learning_rate': 0.1,
    'n_estimators':100
}

xgb_clf = XGB.XGBClassifier(**params)
xgb_clf.fit(x_train_tfidf, y_train)
```

```
XGBClassifier
XGBClassifier(alpha=10, base_score=None, booster=None, callbacks=None,
colsample_bylevel=None, colsample_bynode=None,
colsample_bytree=None, device=None, early_stopping_rounds=None,
enable_categorical=False, eval_metric=None, feature_types=None,
gamma=None, grow_policy=None, importance_type=None,
interaction_constraints=None, learning_rate=0.1, max_bin=None,
max_cat_threshold=None, max_cat_to_onehot=None,
max_delta_step=None, max_depth=4, max_leaves=None,
min_child_weight=None, missing=nan, monotone_constraints=None,
multi_strategy=None, n_estimators=100, n_jobs=None,
num_parallel_tree=None, ...)
```

ACCURACY:  
0.5967824310520939

LOG LOSS XGB:  
0.9029549654440971

The classification report, accuracy and log loss:

	precision	recall	f1-score	support
0	0.53	0.90	0.66	1580
1	0.74	0.38	0.50	1127
2	0.77	0.40	0.53	1209
accuracy			0.60	3916
macro avg	0.68	0.56	0.56	3916
weighted avg	0.66	0.60	0.58	3916

# IMPROVING THE DEEP NEURAL NETWORK ARCHITECTURE:

I Improved the previously used architecture, to a **2 hidden layers** and added a **dropout layer** with **dropout rate 0.3**, below the model architecture and the results obtained after **10 epochs**:

```
Model: "sequential_4"
```

Layer (type)	Output Shape	Param #
dense_12 (Dense)	(None, 128)	1797760
dropout_4 (Dropout)	(None, 128)	0
dense_13 (Dense)	(None, 64)	8256
dense_14 (Dense)	(None, 3)	195

```
=====  
Total params: 1806211 (6.89 MB)  
Trainable params: 1806211 (6.89 MB)  
Non-trainable params: 0 (0.00 Byte)
```

```
Deep Neural Network - Train accuracy: 1.0  
Deep Neural Network - Test accuracy: 0.81
```

	precision	recall	f1-score	support
0	0.80	0.84	0.82	1580
1	0.83	0.80	0.82	1127
2	0.81	0.78	0.80	1209
accuracy			0.81	3916
macro avg	0.81	0.81	0.81	3916
weighted avg	0.81	0.81	0.81	3916

Adding a **dropout layer** and a **batch normalization** layer:

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 128)	1797760
dropout (Dropout)	(None, 128)	0
batch_normalization (Batch Normalization)	(None, 128)	512
dense_1 (Dense)	(None, 64)	8256
dropout_1 (Dropout)	(None, 64)	0
dense_2 (Dense)	(None, 3)	195

=====  
Total params: 1806723 (6.89 MB)  
Trainable params: 1806467 (6.89 MB)  
Non-trainable params: 256 (1.00 KB)

Deep Neural Network - Train accuracy: 1.0  
Deep Neural Network - Test accuracy: 0.8

	precision	recall	f1-score	support
0	0.79	0.82	0.80	1580
1	0.80	0.81	0.81	1127
2	0.81	0.76	0.78	1209
accuracy			0.80	3916
macro avg	0.80	0.80	0.80	3916
weighted avg	0.80	0.80	0.80	3916

## CONCLUSION:

All the depp neural network models and machine learning algorithms gave approxiamte results for this task.